

Lec 5: Linear Regression

Ailin Zhang

Agenda

- Multiple Regression and Least Square
- Maximum Likelihood Estimation
- Gauss Jordan Elimination
- The Sweep Operator

Linear Regression

The dataset of linear regression consists of an $n \times p$ matrix $\mathbf{X} = (x_{ij})$, and a $n \times 1$ vector $\mathbf{Y} = (y_i)$.

$$y_i = \sum_{j=1}^p x_{ij} \beta_j + \epsilon_i,$$

for $i = 1, \dots, n$, where $\epsilon_i \sim N(0, \sigma^2)$ independently for $i = 1, \dots, n$.
Ambiguous about the intercept term.

- $[\mathbf{X}, \mathbf{Y}]$ is called the training data
- y_i is called response variable, outcome, dependent variable.
- x_{ij} is called predictor, regressor, covariate, independent variable, or simple variable.
- In the experimental design setting, \mathbf{X} is called the design matrix.

Linear Regression

obs	$\mathbf{X}_{n \times p}$	$\mathbf{Y}_{n \times 1}$
1	$x_{11}, x_{12}, \dots, x_{1p}$	y_1
2	$x_{21}, x_{22}, \dots, x_{2p}$	y_2
...		
n	$x_{n1}, x_{n2}, \dots, x_{np}$	y_n

- 1 Explanation: understanding the relationship between y_i and $(x_{ij}, j = 1, \dots, p)$.
- 2 Prediction: learn to predict y_i based on $(x_{ij}, j = 1, \dots, p)$, so that in the testing stage, if we are given the predictor variables, we should be able to predict the outcome.

Row Vector treatment

obs	$\mathbf{X}_{n \times p}$	$\mathbf{Y}_{n \times 1}$
1	X_1^\top	y_1
2	X_2^\top	y_2
...		
n	X_n^\top	y_n

$$X_i^\top = (x_{ij}, j = 1, \dots, p)$$

where X_i^\top is the i -th row of \mathbf{X} .

Here X_i is not in bold font.

We can write the model as $y_i = \langle X_i, \beta \rangle + \epsilon_i = X_i^\top \beta + \epsilon_i$, where $\beta = (\beta_j, j = 1, \dots, p)^\top$.

Least Square Method

Least square loss function: $Loss(\beta) = \frac{1}{2} \sum_{i=1}^n \epsilon_i^2$

$\epsilon_i = y_i - s_i$, where $s_i = \sum_{j=1}^p x_{ij}\beta_j$

$$\frac{\partial Loss(\beta)}{\partial \beta_k} = \sum_{i=1}^n \epsilon_i \frac{\partial \epsilon_i}{\partial s_i} \frac{\partial s_i}{\partial \beta_k} = - \sum_{i=1}^n \epsilon_i x_{ik}$$

$$\frac{\partial Loss(\beta)}{\partial \beta_k} = - \sum_{i=1}^n X_i (y_i - X_i^\top \beta) = 0$$

Question: what is the dimensionality of $X_i y_i$ and $X_i X_i^\top$?

Maximum Likelihood

More general than least square $\epsilon_i \sim N(0, \sigma^2)$

$$likelihood(\beta) = \prod_{i=1}^n p(y_i | s_i)$$

Since $y_i = s_i + \epsilon_i$, $[y_i | s_i] \sim N(s_i, \sigma^2)$

Column Vector Treatment

obs	$\mathbf{X}_{n \times p}$	$\mathbf{Y}_{n \times 1}$
1	$\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p$	\mathbf{y}
2		
...		
n		

Geometric Explanation




Figure 1: Least squares projection

Summary: Solve β from $(\mathbf{X}^\top \mathbf{X})\beta = \mathbf{X}^\top \mathbf{Y}$

$$\beta = (\mathbf{X}^\top \mathbf{X})^{-1}(\mathbf{X}^\top \mathbf{Y})$$

Gauss Jordan Elimination - Example

$$\begin{cases} x_1 + x_2 + x_3 = 5 \\ 2x_1 + 3x_2 + 5x_3 = 9 \\ 4x_1 + 5x_3 = 2 \end{cases}$$

Gauss Jordan Elimination

$$\text{GJ}[1 : n][A|b] = [I|A^{-1}b] = A^{-1}[A|b],$$

$$\text{GJ}[1 : n][A|I] = [I|A^{-1}] = A^{-1}[A|I].$$

Gauss Jordan Elimination

For a system of linear equations $Ax = b$

$A = (a_{ij})$ is $n \times n$, $x = (x_i)$ is $n \times 1$, and $b = (b_i)$ is $n \times 1$

we can solve $x = A^{-1}b$ by Gauss-Jordan elimination.

Specifically, for any matrix A (any $n \times N$ matrix)

let $\tilde{A} = \text{GJ}[k]A$, then

$$\begin{aligned}\tilde{A}_k &= A_k / a_{kk}, \\ \tilde{A}_i &= A_i - a_{ik} \tilde{A}_k, \quad i \neq k,\end{aligned}$$

A_k is the k -th row of A . $\tilde{a}_{kk} = 1$, and $\tilde{a}_{ik} = 0$ for $i \neq k$.

- Apply Gauss-Jordan sequentially: $\text{GJ}[1 : m]$ means we apply Gauss-Jordan for $k = 1 : m$.
- Gauss-Jordan is linear: $\tilde{A} = \text{GJ}[k]A \rightarrow \tilde{A} = G_k A$ for a matrix G_k .

R code for Gauss Jordan

```
myGaussJordan <- function(A, m)
{
  n <- dim(A)[1]
  B <- cbind(A, diag(rep(1, n)))
  for (k in 1:m)
  {
    a <- B[k, k]
    for (j in 1:(n*2))
      B[k, j] <- B[k, j]/a
    for (i in 1:n)
      if (i != k)
      {
        a <- B[i, k]
        for (j in 1:(n*2))
          B[i, j] <- B[i, j] - B[k, j]*a;
      }
  }
  return(B)
}

A = matrix(c(1,2,3,7,11,13,17,21,23), 3,3)
myGaussJordan(A,3)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
## [1,]    1    0    0 1.00 -3.0  2.00
## [2,]    0    1    0 -0.85  1.4 -0.65
## [3,]    0    0    1  0.35 -0.4  0.15
```

Computing Efficiency

What is the time complexity?

Can we get rid of any loop?

```
myGaussJordan <- function(A, m)
{
  n <- dim(A)[1]
  B <- cbind(A, diag(rep(1, n)))
  for (k in 1:m)
  {
    a <- B[k, k]
    for (j in 1:(n*2))
      B[k, j] <- B[k, j]/a
    for (i in 1:n)
      if (i != k)
      {
        a <- B[i, k]
        for (j in 1:(n*2))
          B[i, j] <- B[i, j] - B[k, j]*a;
      }
  }
  return(B)
}
```